





Script 1: Training Data Extraction Test

Does DP-QLoRA ($\epsilon \approx 2$) suppress verbatim memorization of training narrations?

Method: model sees first 50% of narration, must complete the rest — suffix never shown



EXTRACTION SUMMARY			
LoRA	train	verbatim	: 0.0%
DP	train	verbatim	: 0.0%
LoRA	val	verbatim	: 0.0%
DP	val	verbatim	: 0.0%
LoRA	train	mean-sim	: 0.1984
DP	train	mean-sim	: 0.1834
LoRA	val	mean-sim	: 0.1699
DP	val	mean-sim	: 0.1629
MAX_LENGTH = 64 (compliance)			
$\epsilon \approx 2.0$ C = 1.5 B = 24			