


Loss Distributions: Members vs Non-Members

More overlap $\rightarrow$ model can't distinguish training data

QLoRA (non-private)
Distribution overlap = 0.872 (higher = better privacy)



DP-QLoRA ($\epsilon \approx 2$, outputs_2)
Distribution overlap = 0.850 (higher = better privacy)

